

LLM-as-Judge Scoring Rubric

NLQ Evaluation Framework — Deliverable 5 (§14.1)

Prompt revision: **judge-87dac1ae3f3b** · generated 2026-09-21

This rubric is rendered directly from the Jinja templates the judge executes, so it cannot drift from the calibrated prompt. Sign-off is against the prompt revision named above.

Purpose. The judge captures answer quality that exact-match cannot measure — completeness, coherence, contextual soundness — without penalising the phrasing, length, tone, or dashboard-styling variation the client has told us to expect (§10, §HC-3). Exact-match accuracy and judge scores are always reported side by side and never blended into a composite (§11.3).

1. The four dimensions

Dimension	What the judge assesses	Anchor scale
Factual correctness	Does the answer match the expected answer in substance? Numeric values must be exact.	5 = exact · 3 = partial · 1 = wrong
Completeness	Does the response include all required values and context from the Q&A; pair?	5 = complete · 3 = missing fields · 1 = severely incomplete
Format adherence	Is the answer in the expected format? Dashboard styling variation is explicitly not a failure.	5 = correct · 3 = minor deviation · 1 = wrong format
SQL plausibility	Is the generated SQL syntactically valid and logically coherent for the question?	5 = valid & correct · 3 = valid but suboptimal · 1 = invalid

2. Full 1–5 scale (levels 2 and 4 require Platform Owner sign-off)

Score anchors 5, 3 and 1 are verbatim from the Execution Scope. Levels 2 and 4 are interpolated by Codehive so the judge has a continuous scale to reason toward. **This interpolation is the item requiring sign-off on this document.**

FACTUAL CORRECTNESS - does the answer match the ground-truth expected answer in substance?
5 = exact: the stated fact matches the expected answer, wording may differ freely
4 = correct but with a minor imprecision (rounding, units restated loosely)
3 = partial: some stated facts correct, others wrong or unsupported
2 = mostly wrong with an incidental correct element
1 = wrong: the central fact contradicts the expected answer

COMPLETENESS - does the response include all required values, columns, or rows?
5 = complete: everything the ideal response carries is present, including the correct entities and ordering when the question asks for them
4 = all key values present, but a minor supporting detail is omitted
3 = a required value, label, entity, or row is missing or wrong
2 = only a fraction of what was asked for
1 = severely incomplete: the question is essentially unanswered

FORMAT ADHERENCE - is the answer in the expected format (number, table, list, sentence)?
5 = correct: presentation matches what the question calls for
4 = correct format with cosmetic deviation
3 = minor deviation: readable but not the shape asked for
2 = wrong shape, content still recoverable
1 = wrong format: presentation obstructs the answer

SQL PLAUSIBILITY - is the generated SQL syntactically valid and logically coherent?
5 = valid & correct: would produce the expected answer
4 = valid and correct, stylistically different from the reference
3 = valid but suboptimal: works, but joins/filters are loose or redundant
2 = valid syntax, logically wrong for the question
1 = invalid: will not parse or execute

3. Non-determinism boundary and worked examples (§HC-3, §10.2)

Encoded verbatim in every judge prompt. The worked examples are load-bearing: without them the model over-penalises phrasing variation, which corrupts every downstream regression comparison.

NON-DETERMINISM BOUNDARY – READ CAREFULLY (§HC-3)

Numeric values MUST be identical to the EXPECTED ANSWER. Any numeric variance is a factual failure. Wording, sentence length, ordering of prose, tone, currency formatting, and dashboard presentation are NEVER deductions. Only substance and numeric accuracy are scored.

WORKED EXAMPLES – these are how the boundary applies in practice.

DO NOT DEDUCT (all four should score 5 on the dimensions shown):

Example A – Prose wrapper around the number, no factual change.

EXPECTED ANSWER: 72

PLATFORM ANSWER: "There are 72 active accounts in the CRM."

→ factual_correctness = 5, completeness = 5, format_adherence = 5.

The prose wrapper is not a deduction. The number matches exactly.

Example B – Currency formatting differs, number identical.

EXPECTED ANSWER: Priya Raghavan – 4,182,650.00

PLATFORM ANSWER: "Priya Raghavan closed \$4,182,650 in Q2, a strong quarter."

→ factual_correctness = 5, completeness = 5.

Currency prefix and dropped `.00` are cosmetic. Rep name and number match.

Example C – Different wording, same substance.

EXPECTED ANSWER: 5 orphan contacts

PLATFORM ANSWER: "Five contacts have no account assigned."

→ factual_correctness = 5. Spelled numeral is not a factual defect.

Example D – Sentence order and length differ, same information carried.

EXPECTED ANSWER: Technology 8,240,100.00; Finance 6,102,750.00; Healthcare 4,988,300.00

PLATFORM ANSWER: "H1 leaders were Healthcare (4,988,300.00), Finance (6,102,750.00), and Technology (8,240,100.00)."

→ factual_correctness = 5, completeness = 5. Order reversed but every value is present with identical numerics. Not a deduction.

DO DEDUCT – the score to assign is stated on each example:

Example E – Number wrong, wording irrelevant.

EXPECTED ANSWER: 72

PLATFORM ANSWER: "There are 68 active accounts."

→ factual_correctness = 1. Any numeric variance is zero tolerance.

Example F – Number rounded to fewer digits than expected.

EXPECTED ANSWER: 4,182,650.00

PLATFORM ANSWER: "4,182,000"

→ factual_correctness = 1. Rounding-loss is still numeric variance.

Example G – Required entity missing, value present.

QUESTION: "Which sales rep closed the most deal value, and how much?"

EXPECTED ANSWER: Priya Raghavan – 4,182,650.00

PLATFORM ANSWER: "4,182,650.00"

→ factual_correctness = 5, completeness = 3. The number is right, so nothing factual is contradicted. One of the two things asked for is missing, which is "missing fields" on the completeness scale – a 3, NOT a 1. Reserve 1 for a response that leaves the question essentially unanswered.

Example H – Answer is entirely verbal, no numbers where numbers were asked.

EXPECTED ANSWER: Technology 8,240,100.00; Finance 6,102,750.00; Healthcare 4,988,300.00

PLATFORM ANSWER: "Technology led H1 2026 in closed deal value."

→ factual_correctness = 1, completeness = 1. The specific numbers were the substance of the answer; a category without figures does not answer this question, so here completeness genuinely is 1.

RULE OF THUMB: if you find yourself deducting for HOW something is said rather than WHAT is said, stop and re-check against the examples above.

4. Per-dimension scoped instructions

In *per_dimension* mode each dimension is scored in its own prompt to remove the halo effect. These are the scoped instruction blocks.

4.1 Factual Correctness

Assess ONLY whether the Platform answer matches the ground-truth EXPECTED ANSWER in substance. Equivalent wording, different sentence structure, spelled-out numerals, and added units are all acceptable and must not reduce the score.

5 = exact: the stated fact matches the expected answer, wording may differ freely
4 = correct but with a minor imprecision (rounding, units restated loosely)
3 = partial: some stated facts correct, others wrong or unsupported
2 = mostly wrong with an incidental correct element
1 = wrong: the central fact contradicts the expected answer

Ignore presentation, ignore omissions of secondary detail, ignore the SQL.

CRITICAL: this dimension penalizes ONLY facts that CONTRADICT the expected answer. A vague, evasive, or incomplete answer that does not assert anything false is NOT a 1. If the platform names a correct entity or reports a correct value – even if it omits others – that IS partial factual content and must be scored 3, 4, or 5, never 1 or 2. Score 1 or 2 only when a stated fact is positively wrong (wrong number, wrong entity, wrong direction). Missing content is a completeness issue, not a factual one.

4.2 Completeness

Assess ONLY whether the Platform answer includes all required values, columns, rows, and required named entities from the JUDGE REFERENCE. The response is incomplete if it reports the right numeric shape but the wrong entity label (for example, the wrong account, region, rep, industry, or ranking order). Use the JUDGE REFERENCE as the definition of "complete" - it is the ideal full response.

5 = complete: everything the ideal response carries is present, including the correct entities and ordering when the question asks for them
4 = all key values present, but a minor supporting detail is omitted
3 = a required value, label, entity, or row is missing or wrong
2 = only a fraction of what was asked for is present
1 = severely incomplete: the question is essentially unanswered

Do not reward extra length or unrequested commentary - verbosity is not completeness. Ignore factual accuracy and presentation style here, except when the wrong entity or wrong order means the response no longer delivers the full requested content.

CRITICAL: a bare value without its entity label is NOT complete. Examples:

- Question "How many active accounts do we have?" with answer "72" – the value is present but the required entity label ("active accounts") is missing. Score 3, not 5.
- Question "Which region had the most closed-won opportunities?" with answer just "West" (missing count) – a required value is missing. Score 3, not 5.

An answer that emits every required VALUE but omits required LABELS or QUALIFIERS is a 3 per this dimension. Only bare values that fully saturate the question (e.g. "5" as the answer to "how many contacts have no account?" when the question itself supplies the entity context) can reach a 4 or 5.

4.3 Format Adherence

Assess ONLY whether the answer is presented in a form appropriate to the question – a single number, a table, a list, or a natural-language sentence – using the JUDGE REFERENCE as the model of the expected shape.

5 = correct: presentation matches what the question calls for
4 = correct format with cosmetic deviation
3 = minor deviation: readable but not the shape asked for
2 = wrong shape, content still recoverable
1 = wrong format: presentation obstructs the answer

Minor formatting differences must NOT be heavily penalised. Ignore whether the content is factually right – that is scored elsewhere.

4.4 SQL Plausibility

Assess ONLY the SQL the Platform generated: is it syntactically valid and logically coherent for the question? The ground-truth reference SQL is not available to you; assess only the question and the Platform-generated SQL.

- 5 = valid & correct: would produce the expected answer
- 4 = valid and correct, stylistically different from the reference
- 3 = valid but suboptimal: works, but joins/filters are loose or redundant
- 2 = valid syntax, logically wrong for the question
- 1 = invalid: will not parse or execute

A valid formulation using different aliases, join order, or aggregate spelling is not a defect. You are judging query quality, not execution performance. If the Platform returned no SQL at all, score 1.

5. Determinism, structured output and audit settings

- Temperature is fixed at 0 and a fixed seed is sent wherever the provider supports it (§10.1) — a non-deterministic judge cannot certify determinism. Some deployments (GPT-5 / o1 family) refuse temperature 0 outright; by default that fails the run, and the operator may instead opt in to running on the model default plus the fixed seed. A run that does so records `judge_temperature_enforced=false` on its scorecard, so the determinism basis of every run is visible on the artefact. Anthropic has no seed parameter at all, so a run on an Anthropic model records `judge_seed_enforced=false` rather than implying a control it never applied.
 - The judge receives the question, the platform response, the `expected_answer` and the `judge_reference`. It never receives the `reference_sql` — SQL plausibility is scored on the platform's own generated SQL (§10.1).
 - Structured output is schema-validated: four integer scores (1–5, never coerced from strings or floats) plus one rationale per dimension. Malformed output is retried, then failed loudly (§10.1).
 - The full prompt and full raw response are logged for every scored item under a stable `judge_run_id`, including cache hits (§10.1).
 - Calibration must pass before any production scoring run: within ± 1 of the human anchor score on $\geq 90\%$ of anchors per dimension, with no directional disagreement, over ≥ 10 anchors per domain (§10.2). Uncalibrated scores never enter a scorecard (§10.2). A pass is recorded against the judge that earned it — model, prompt revision and scoring mode — and a run differing on any of the three is treated as uncalibrated.
 - Anchor sets are themselves checked: §10.2 asks for anchors "spanning the score range — not 10 easy passes", so a set on which a constant score would satisfy the acceptance test is rejected before any provider budget is spent on it.
-

6. Deterministic comparison — decisions requiring sign-off

Exact-match is not part of the judge, and the two are never blended (§11.3). It is documented here because OI-2 and OI-3 are still open, and the delivery team has settled both in order to produce an interpretable accuracy number. **These four decisions are the second item requiring sign-off on this document.**

- **Numerics compare by value, not by rendering.** SHC-3 protects phrasing, length and presentation and fails only numeric variance, so a thousands separator, a currency prefix and a trailing `.00` are not variance — 28,731 equals 28731. A rounded figure still fails: 4,182,000 against 4,182,650.00 is a real loss of value. **This is Codehive's reading of OI-2**, taken so the scorer is usable now, and it is one flag to reverse.
 - **Entity labels are required, and each must be attached to its own value.** An expected answer of Priya Raghavan — 4,182,650.00 is not satisfied by Bob Smith closed 4,182,650.00, nor by a list whose values are each bound to the wrong category. Comparing numerics alone passed both.
 - **Dates compare by day, with no tolerance.** Material change C10 withdrew the ± 1 day tolerance carried in earlier scoring drafts. ISO on the ground-truth side; ISO or a month-name rendering on the platform side, because the platform writes dates in prose. **This is Codehive's reading of OI-3.**
 - **Spelled-out numerals are deliberately not recognised by the scorer.** The judge scores "Five contacts" against an expected 5 as a 5, and should. Teaching the deterministic scorer number words would make an expected 1 match any answer containing the word "one" — a common false pass traded for a rare false fail. The two scores diverging on such an answer is intended, and both are always shown (§11.3).
 - A run that relaxes any of the above is marked on its scorecard as `comparison_policy` and cannot establish a baseline.
-

Sign-off

Approved for calibration and production scoring against prompt revision **judge-87dac1ae3f3b**, including the interpolated score levels 2 and 4 (section 2) and the deterministic comparison decisions in section 6.

Platform Owner	_____	Date	_____
Codehive lead	_____	Date	_____