

Regression scorecard ▪ PREVIEW

domain	exact-match	of eligible	vs 95% gate	judge overall	vs baseline
crm	55.00%	88 of 160	BELOW GATE	3.59	no baseline

Read first

- **ACCURACY GATE:** crm is at 55.00% against the 95% required before the package ships (§14.2 condition 1).
- **PREVIEW run** — not an official evaluation. It neither establishes nor updates a baseline and cannot certify a release (§10.2, §11.3).
- The judge is **UNCALIBRATED**. These scores are diagnostic only and must not feed a release decision (§10.2).
- **Assembled:** combined from crm/merged-20260921T124500Z-20260921T153322Z-RECONSTRUCTED


Run provenance

run_id	combined-20260923T131009Z
run_timestamp_iso	2026-09-23T13:10:09.525769+00:00
platform_version	not supplied
dataset_version	not supplied
scorecard_mode	PREVIEW
judge calibrated	NO
exact-match comparison	numeric=value labels=required proximity=120
determinism	seed never reached provider

Accuracy gate (§14.2) — all nine before the package ships

§14.2 condition	status	basis
1. >=95% exact-match, every domain	FAIL	below the gate: crm
2. Numeric values identical on repeat runs	AT RISK	exact-match is deterministic; the judge's own determinism is NOT fully enforced this run
3. Independent reviewer re-executes every reference_sql	MANUAL	§15 protocol — not observable from a run
4. Failing pairs reworked, never removed	PASS	no questions excluded from the denominator
5. Generated SQL logged for every failure	PASS	platform_generated_sql is written for every question in question_results.csv
6. Cross-domain question uniqueness	MANUAL	§15 normalise-and-hash across the corpus
7. Baseline run completed across all domains	NOT YET	this is a PREVIEW scorecard; only a RELEASE run establishes a baseline
8. Clean-room re-run reproduces byte-identical CSVs	MANUAL	asserted by the dataset pipeline, not by this run
9. CRM certified before final delivery	FAIL	crm is below the gate

Exact-match by tier — crm



Pair counts against the §9.1 quota

domain	tier	questions	§9.1 quota	
crm	T1	32	32	ok
crm	T2	40	40	ok
crm	T3	32	32	ok
crm	T4	32	32	ok
crm	T5	24	24	ok

A tier short of quota means a pair left the set. §14.2 condition 4: failing pairs are reworked, never removed — removing one raises the percentage without improving anything.

Where the failures are

count	failure kind — 72 failure(s) in total
68	value <value> not present
4	label <value> present but a different label claims its value

Grouped by kind — quoted and numeric literals are collapsed, so one finding seen many times reads as one row. Every failure's full detail and its platform-generated SQL are in question_results.csv (§14.2 condition 5).

Deterministic accuracy (§HC-3, zero numeric tolerance)

domain	tier	questions	pass	exact-match %	n/a	clarified	baseline %	delta pp	regression
crm	ALL	160	88	55.00	0	0	—	—	—
crm	T1	32	26	81.25	0	0	—	—	—
crm	T2	40	30	75.00	0	0	—	—	—
crm	T3	32	32	100.00	0	0	—	—	—
crm	T4	32	0	0.00	0	0	—	—	—
crm	T5	24	0	0.00	0	0	—	—	—

Judge quality (1–5 per dimension, §10)

domain	tier	factual	complete	format	sql	overall	clarified	judge err	platform err	null-handling
crm	ALL	3.34	3.64	3.58	3.81	3.59	0	0	0	0
crm	T1	4.25	4.31	3.97	4.38	4.23	0	0	0	0
crm	T2	3.98	4.15	3.85	4.38	4.09	0	0	0	0
crm	T3	5.00	5.00	4.41	5.00	4.85	0	0	0	0
crm	T4	1.09	1.94	2.84	2.22	2.02	0	0	0	0
crm	T5	1.88	2.38	2.50	2.62	2.34	0	0	0	0

Exact-match and judge scores are reported side by side and never combined into a composite (§11.3) — they measure different things. Baseline comparison and the regression flag are domain-level; tier rows are diagnostic. **n/a** counts pairs with no deterministic core and **clarified** counts questions the platform declined to answer; both sit outside the percentage, so a rise in either is not a rise in accuracy (§14.2 condition 4). **null-handling** is the §9.2 T3 diagnostic and is never part of either score.